

CLOSING SESSION

Closing: what we learned, what we take home

Moderator: Vivek · Convene Philadelphia · Oct 14, 2026

PANEL OPENER

When AI surprised you

Describe a moment when an AI tool in your function surprised you, did something you didn't expect, or worked around a limit.

What was your human-check setup at the time, and what did you change afterward?

PANEL · Q1

Is the HCP relationship shifting?

Have you felt the HCP relationship shifting in your function: people getting answers from AI assistants instead of your MSLs, your portal, your team?

What's the most concrete example, and how did you respond?

PANEL · Q2

What should an agent do, and never touch?

Coinbase's fraud agent can only score risk; it can't act, and it can't override the human system. In your function:

DO on its own	DRAFT FOR SIGN-OFF a human approves	NEVER TOUCH humans only
Which tasks should an agent do on its own?	**Which should it only draft for human sign-off?**	**Which should it never touch?**

PANEL · Q3

Safe-to-fail experiments

The research says the winners just kept experimenting; most successful AI projects had failures nobody was punished for.

Tell us about a safe-to-fail experiment you ran in your function. What made it safe, and what did you learn?

PANEL · Q4

Train experts to challenge AI

The data says the biggest AI gains go to people who are best at catching the machine's mistakes, not those who trust it most.

How are you training your experts (MSLs, reviewers, analysts) to challenge AI output instead of deferring to it? Give us something concrete that's worked.

PANEL · Q5

Design the exam for an AI agent

Legal now grades its AI agents against strict pass/fail tests judged by real practitioners.

If you had to design the exam for an AI agent in your function, say one answering HCP questions, what would the test be, and who would you trust to grade it?

PANEL · Q6

Taking a pilot to legal / compliance

Think of the last time you took an AI pilot to legal or compliance.

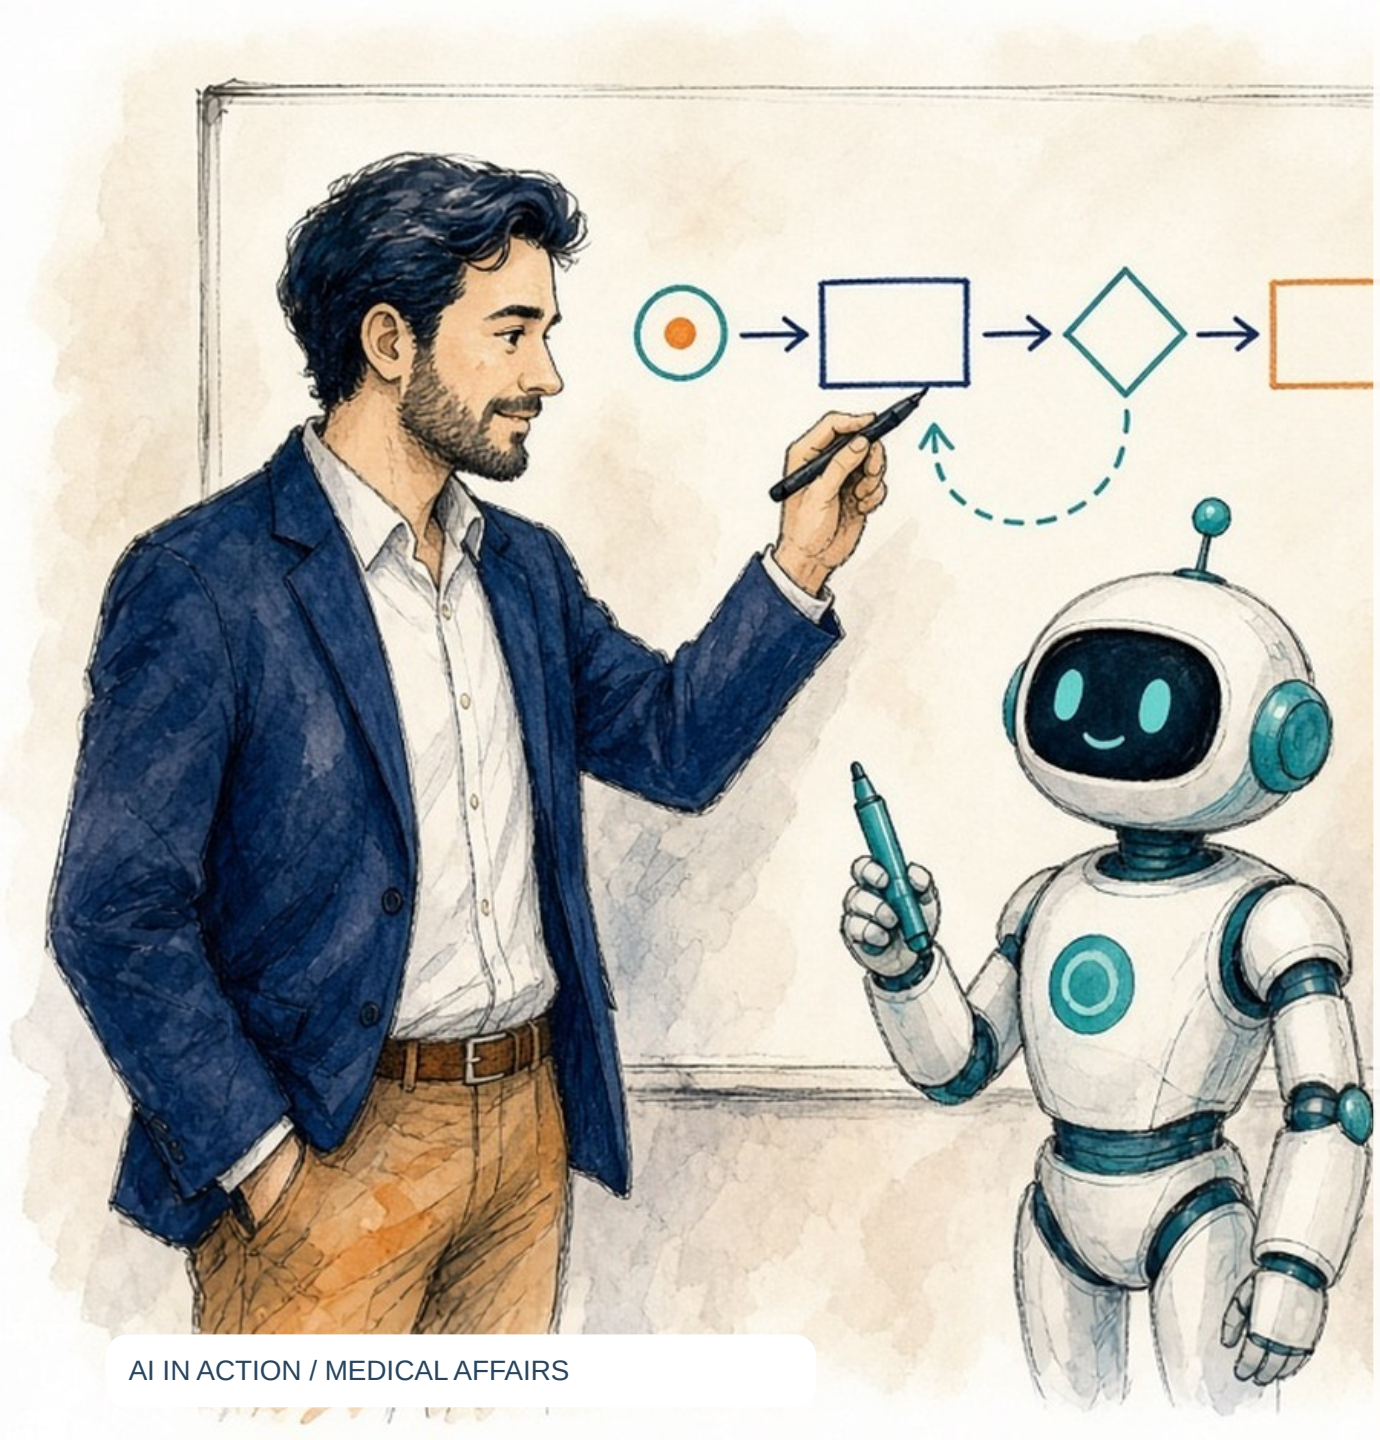
How did you frame the risk, and looking back, did you get the calibration right? What would you do differently?

PANEL · Q7

From hackathon to real pilot

Two days ago most of this room had never built an agent; yesterday you demoed them.

Name the one thing from your hackathon team going into a real pilot, and the one thing that must be true before it touches a real HCP.



YOUR TURN

Your turn: audience

Short answers from the room.
Hands up when called.

1. What's the one thing from these two days you'll actually use in your work next week?

2. For the MSL and med-info folks: what's the task you'd most want an agent to take off your plate, and the one you'd never hand over?

3. What was the moment in these two days when you went from skeptical to 'okay, this could actually work'? What triggered it?

4 What's something you believed about AI in medical affairs on Monday morning that you no longer believe?

5 Who here tried to break someone else's agent in the red-team swap, and what did breaking it teach you about your own build?

6 Think back to the workflow you mapped on day one. After two days of building, what's the part you now realize is harder to automate than you expected?

7 What's the one prompt, routine, or agent behavior going onto your deployment commitment form, and what's the first thing you'll do with it on Monday?

8 For the leaders in the room: what conversation do you need to have back home before any of this becomes real, and with whom?

9 Show of hands: who arrived worried AI would replace part of their function, and who leaves thinking it will change the function instead? What changed your mind?

Human at the helm

Opportunity AI

Context

Trust

ACROSS MEDICAL AFFAIRS

MSL pre-call	Medical information
Congress intelligence	Publications
Field insights	Advisory boards
Launch planning	Evidence generation

Medical Affairs is the steward of trust. AI helps us keep it.

