




02

Insights Engine

Turn a quarter of field records into decisions leadership can act on.

Hackathon vertical 02 of 04 · build afternoon agenda, swarm, gates, question bank and impact worksheet

01 Publications Brain

02 Insights Engine

03 Congress Monitor

04 Field Intelligence Engine

AI in Action for Medical Affairs

October 13–14, 2026 · Convene Philadelphia (2nd floor), Two Commerce Square, 2001 Market St

Host and keynote: Vivek Mukhatyar · aiinaction.up.railway.app/hackathon



Scan for the hackathon pack

The vertical on one page

THE CHALLENGE

Every quarter MSLs log hundreds of field observations and sit in advisory boards. By the time someone codes, themes and writes them up, the quarter is over, and leadership gets theme counts instead of insights it can act on.

DAY 2 SLOT · WED OCT 14

10:00–10:20 AM

About 20 minutes, give or take, including a few questions. Teams present 01 → 04 from 9:40.

THE PROBLEM PROMPT

Build an AI workforce that reads a quarter of field records (and a meeting transcript), surfaces possible safety findings first, and turns the rest into the three insights leadership most needs: what is happening, why it matters, what it changes, with source record IDs, a stated confidence, an owner and a next action for each, while keeping disagreements, weak signals and what the field did not say. Humans approve the insights and the brief.

Teams can keep this prompt or narrow it to one workflow they know. A good narrowed prompt names the role, the trigger, the deliverable and the decision it serves.

DONE LOOKS LIKE · DAY 1

The swarm ran end to end at least once on practice data; it stopped at Gate 1 and Gate 2 and the team made both decisions; screenshots and minute-level timings are in the capture log; the agent has drafted a slide outline.

DONE LOOKS LIKE · DAY 2

The deck is built from the final template: no [brackets] left, impact numbers are the team's own, practice data marked fictional, the LIVE DEMO slide left empty, backup screenshots in place. The team presents at 10:00 for about 20 minutes with a short live demo.

LEAD SKILL

field-insight-synthesis

SUPPORTING SKILLS

insight-generation

strategic-analysis

meeting-transcription

DATA PACK

Any TA. Each field-observations.csv contains deliberately seeded adverse events, a product quality complaint and an off-label use: the safety-first stop will trigger.

field-observations.csv advisory-board-transcript.md

interaction-notes.md medical-plan.md

evidence-landscape.md + practice CRM

Fictional practice data (Nordvant Biopharma), or bring your own non-confidential data.

WARM-UP MISSIONS

L2 field-insights

What should leadership know?

L2 transcript

Turn a meeting transcript into action



The build afternoon, block by block

A suggested, optional flow: **Grok Bot setup (30 min)** → **warm-up missions** → **workflow ideation** → **the hack**. Teams may move faster or slower. **5:00 PM: no formal share-out**; facilitators circulate and ask questions (page 6), then the cocktail reception.

1 · Setup 30 min	2 · Warm-ups 40 min	3 · Ideation 45 min	4 · The hack 80 min
1:45	2:15	2:55	3:40
5:00			
1 Set up Grok Bot + skills 1:45–2:15 · 30 min	2 Warm-up missions 2:15–2:55 · 40 min	3 Workflow ideation 2:55–3:40 · 45 min	4 The hack 3:40–5:00 · 80 min
1:45 Huddle 1:50 Everyone opens aination.up.railway.app/grokb 2:02 The driver pastes the setup prompt below 2:10 Smoke test passes	2:15 Split into two pairs 2:18 Run both missions in parallel 2:40 Push on it 2:50 Swap outputs	2:55 Map today's workflow for one real quarterly field- insight... 3:10 Spot the AI ideas 3:22 Pick one problem 3:32 Sketch the swarm on paper	3:40 Write the assignment 3:50 Run 1 4:25 Optional stretch break 4:30 Run 2 4:50 Capture
DONE LOOKS LIKE Every member signed in; the driver's conversation has the skills loaded and has read one practice file; capture log started	DONE LOOKS LIKE A leadership brief and a decision log exist; the seeded safety finding was raised first and read aloud; 3 observations logged	DONE LOOKS LIKE The quarterly insight cycle mapped; one problem chosen; hours per quarter and number of people involved estimated; ideas marked Efficiency or Opportunity	DONE LOOKS LIKE The swarm produced an insight table with source IDs and confidence, then a brief, stopping at both gates; the safety list came first; screenshots and timings captured
FACILITATORS CIRCULATE Check every laptop is signed in and the smoke test shows a file marked SYNTHETIC. GitHub blocked? Point to the starter bundle on /grokb.	FACILITATORS CIRCULATE Let the missions finish without steering. Watch for the safety stop and the human-review moment; keep the 2- minute swap on time.	FACILITATORS CIRCULATE Ask the ideation questions (page 6). Push for one problem statement and a baseline estimate before they sketch the swarm.	FACILITATORS CIRCULATE Ask the hack, gate and data questions. From 4:45, check that screenshots and the capture log are saved.

DAY 2 · WED OCT 14 · FINAL BUILD → THIS TEAM'S SLOT

9:00–9:38 · Final build + rehearsal re-open → last fix, agent fills the template → team review → rehearse · MC opens 9:38	9:40 01 Publications Brain	10:00 02 Insights Engine	10:20 03 Congress Monitor	10:40 04 Field Intelligence Engine	11:00 WPP demo
---	---	---------------------------------------	--	--	-----------------------------

Insights Engine: 10:00–10:20 AM, about 20 minutes, give or take, including a short live demo and a few questions. Teams present 01 → 04 from 9:40. Keep all decks on one laptop and treat the Day 1 4:50–5:00 capture as essential.



Warm up, then pick your data

Warm-up missions Block 2 · 2:15–2:55

Level 2 · Pair

Pair A

“What should leadership know?”

The core insight task. Watch for the safety finding the agent should raise before any analysis.

aiinaction.up.railway.app/missions/field-insights

Level 2 · Pair

Pair B

“Turn a meeting transcript into action”

Advisory-board transcript to reviewed transcript, correction queue and decision log: a second insight source.

aiinaction.up.railway.app/missions/transcript

STRETCH, ALTERNATIVES AND REFERENCES

L1 Starter

L2 Pair

L3 Team

L4 Swarm

L1 Alternative

Turn the call notes into useful follow-up

</missions/msl-post-call>

L3 Hack reference

Work across a practice CRM and content library

</missions/connected-planning>

L3 Optional stretch

Run the next 30 days of Medical Affairs

</missions/thirty-day-capstone>

L4 Also useful

Build the whole medical launch plan

</missions/launch-plan-swarm>

BG Background

Team mission 2: the original team mission

[Medical-Affairs-Skills · workshop](#)

Data pack browse and copy links at aiinaction.up.railway.app/data

Fictional practice data, or bring your own non-confidential data.

Nordvant Biopharma and its products are fictional. Own data: no patient data, no confidential company data, no real HCP personal data.

PICK ONE TA

MM · NORVANTIB

AD · DERMALYX

Obesity · ADIPOSYN

Any TA. Each field-observations.csv contains deliberately seeded adverse events, a product quality complaint and an off-label use: the safety-first stop will trigger.

TA PACK FILES synthetic/<TA>/

TA field-observations.csv

TA advisory-board-transcript.md

TA interaction-notes.md

TA medical-plan.md

TA evidence-landscape.md

TA product-profile.md

TA terminology-coding-queue.csv

TA medical-impact-metrics.csv

PRACTICE CRM synthetic/connected/

CRM interactions.csv

CRM hcps.csv

CRM accounts.csv

CRM engagement.csv

CRM data-dictionary.json

The CRM interactions table links every row back to its source file and record ID (e.g. field-observations.csv, OBS-001): ask the swarm to keep those IDs in every insight.



Efficiency vs Opportunity, and the digital team

Efficiency AI today's work, faster

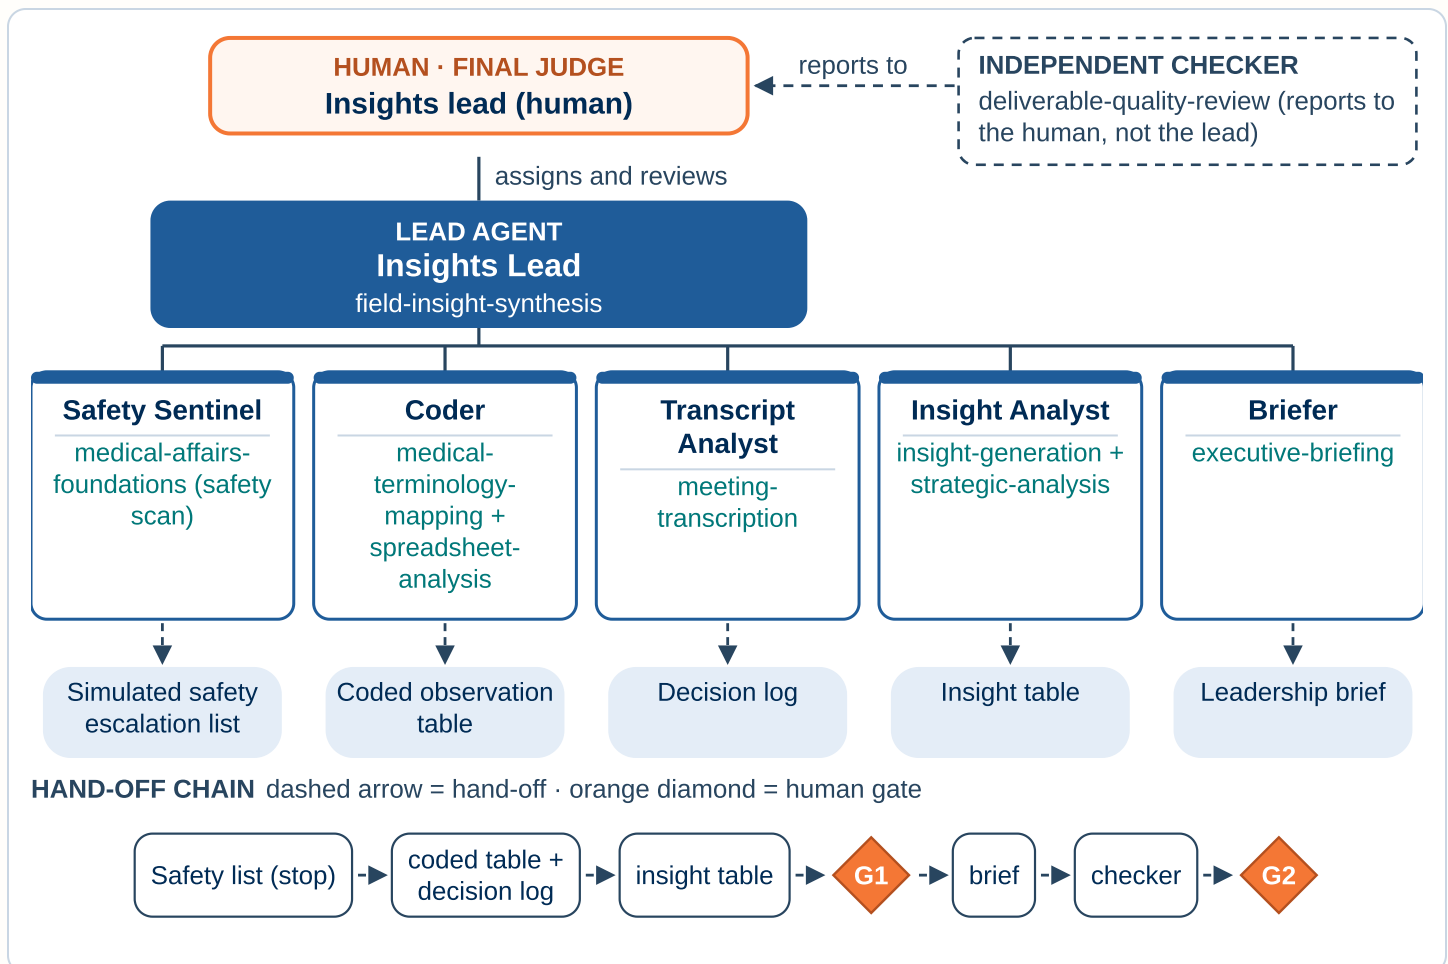
- E1** Auto-coding: every observation coded to themes and controlled terms, de-duplicated, with source IDs, in minutes instead of days.
- E2** First draft of the quarterly insight report (insight, evidence, confidence, owner, action) ready for the insights lead to edit.
- E3** Advisory-board transcript to a decision and follow-up log the same afternoon.

vs

Opportunity AI something not possible today

- O1** A weekly insight pulse instead of a quarterly report: leadership hears about a shift while it can still act.
- O2** Plan-assumption watch: link every insight to the medical plan objective it supports or contradicts, and flag the assumptions the field is challenging.
- O3** Weak-signal and silence detection across all three TAs: the things only one or two MSLs mentioned, and the topics the field conspicuously did not raise.

Suggested swarm one lead agent, sub-agents with their skills, two human gates



Where humans decide, and what to ask



G1

Gate 1

WHEN After the Insight Analyst

WHO DECIDES The insights lead

WHAT THEY CHECK Approves the top three insights, their confidence and the safety routing before the brief is written.

WHY A HUMAN What counts as an insight is a judgment, and safety routing is a regulated obligation.

G2

Gate 2

WHEN After the Briefer

WHO DECIDES The head of medical

WHAT THEY CHECK Signs off the brief, the owners and the actions before it is shared beyond the team.

WHY A HUMAN Accountability for decisions and for what leadership is told.

Always-on stop (a rule, not a gate): any possible adverse event, product quality complaint or off-label signal in human-sourced records is listed verbatim and routed by a human before analysis continues.

Facilitator question bank the Day 1 feedback loop: ask one or two per visit, then move on

Ideation 2:55–3:40

- 1 What is the difference between an observation and an insight in your team today? Who decides?
- 2 How long after the quarter ends does leadership see the insights today? What decision is late as a result?
- 3 Which data would you need from the CRM that the field notes do not have?
- 4 What would leadership do differently with a weekly pulse instead of a quarterly report?

The hack 3:40–5:00

- 5 How does the swarm state confidence, and what would lower it?

Gates during the hack

- 6 Where is Gate 1, and what exactly does the insights lead see there?

Data during the hack

- 7 Did your swarm raise the seeded safety finding before anything else? Who receives it, and what does the agent stop doing?
- 8 Show me one insight and the record IDs behind it. How many records support it?
- 9 Where do disagreements between MSLs end up? Did any get smoothed over?

Impact during the hack

- 10 How many hours does one quarterly cycle take today, across everyone involved? What does the AI-assisted cycle take, including review?



When a team gets stuck

The ten snags facilitators see most often in this vertical, each with the unblocker to offer. Nudge, then step back: the team does the building.

Pitfall	How to unblock
01 Themes with counts, not insights.	Ask: "For each, what is happening, why it matters, what it changes and who should act?" Load insight-generation.
02 The safety scan is skipped or buried.	Make the Safety Sentinel the first worker and tell the lead to stop until a human acknowledges the list.
03 Over-confident insights from one or two records.	Require a confidence level and the record count for every insight; weak signals go in their own section.
04 Source IDs disappear in the brief.	Tell the Briefer to keep IDs in an appendix table; the checker fails any insight without them.
05 Public literature mixed into fictional data.	Keep public evidence in a separate, labelled section, or leave it out.
06 Too much data to read in the time.	Start with one TA and one quarter; add the CRM tables in the second run.
07 Grok Bot cannot open GitHub.	Upload the starter bundle for your TA from https://aiinaction.up.railway.app/grokbob, or the repository ZIP.
08 The agent asks to connect company systems.	Say no: the workshop uses practice data. It is in every setup prompt.
09 A run is slow or wanders.	Ask for the plan first, then run one sub-agent at a time; cut the scope to one TA, one HCP, one output.
10 The output is generic.	Teach the agent: add 3–5 house rules (what good looks like, what must never appear) and run again.



Unblock, then step back. Teams learn most when they make the call at each gate themselves.



Impact worksheet and setup prompt

Impact worksheet the team's own estimates

One cycle = one quarterly field-insight report

Today baseline	h	×		×	\$ /h	=	\$
	hours per cycle (all people)				loaded rate		cost per year
vs							
AI-assisted with the swarm	h	×		×	\$ /h	=	\$
	hours per cycle, incl. human review				loaded rate		cost per year
HOURS SAVED PER YEAR (today h – AI h) × cycles =		VALUE PER YEAR (ESTIMATE) hours saved × rate =		OTHER BENEFITS			
h		\$		days speed to decision	% coverage	# errors caught	

ILLUSTRATIVE

Not data, only an example: 40 h → 8 h per cycle × 4 cycles/year = 128 h saved; × \$150/h = \$19,200 per year. Replace every number with the team's own estimate.

Count every person involved and include human review time in the AI-assisted row. These numbers go on the Impact slide of the final template (slide 9: formula tiles and a bar chart built from the team's own estimates).

Setup prompt Block 1 · 2:02 · the driver pastes it into Grok Bot (replace the TA)

1 · Setup prompt

[copy all](#)

Use <https://github.com/Open-Medical-Affairs/Medical-Affairs-Skills>
 Read AGENTS.md. We are the Insights Engine team at AI in Action for Medical Affairs.
 Load medical-affairs-foundations and these skills: field-insight-synthesis, insight-generation, meeting-transcription, medical-terminology-mapping, executive-briefing, deliverable-quality-review.
 Use the practice data in <https://github.com/Open-Medical-Affairs/Data-Sources>: synthetic/[oncology-mm | immunology-ad | cardiometabolic-obesity]/ and synthetic/connected/ (the practice CRM). It is fictional (Nordvant Biopharma).
 First, only: list the skills you loaded, then show the first 5 rows of field-observations.csv and confirm the file is labelled SYNTHETIC. Do not start a mission yet.
 Do not ask me to connect company systems. Do not send, post or change anything outside this conversation.

SUGGESTED ROLES (ROTATE IF YOU LIKE)

Driver

types into Grok Bot

Navigator

reads the mission and
pushes back

Scribe

owns the capture log

Timekeeper + presenter

keeps the clock, tells the
story



Assignment starter for the hack

Paste into Grok Bot at the start of Block 4 (3:40), then edit the [brackets]. The setup prompt for Block 1 is on page 8.
Ready prompts for every mission: aiinaction.up.railway.app/prompts · swarm mode and the Skill creator: [/optimizer](#).

2 · Assignment starter: Insights Engine swarm

Block 4 · 3:40

```
# Assignment: Insights Engine swarm
You are the Insights Lead, a lead agent coordinating a small team of digital workers for a Medical Affairs team. This is an assignment, not a question: plan, delegate, check and hand back finished drafts.

## Goal: what done looks like
[One sentence: the problem we picked in ideation, and the deliverable that solves it.]

## Context to use
- https://github.com/Open-Medical-Affairs/Medical-Affairs-Skills : read AGENTS.md, load medical-affairs-foundations, your lead skill field-insight-synthesis, and the skills named below.
- Practice data (fictional): https://github.com/Open-Medical-Affairs/Data-Sources synthetic/[TA]/ and synthetic/connected/. Or our own non-confidential data: [describe].

## How to work
1. Show me a short plan and the org chart of workers first. Wait for my OK.
2. Scan human-sourced records for possible safety findings (adverse events, product complaints, off-label use) before anything else. List them verbatim and stop until I acknowledge.
3. Run these sub-agents, each with only the context it needs:
  - Safety Sentinel (skills: medical-affairs-foundations (safety scan)): Reads every record first; lists possible AEs, PQCs and off-label mentions verbatim and stops for routing. Hands off: Simulated safety escalation list.
  - Coder (skills: medical-terminology-mapping + spreadsheet-analysis): De-duplicates contacts, codes observations, keeps source IDs. Hands off: Coded observation table.
  - Transcript Analyst (skills: meeting-transcription): Turns the advisory-board transcript into decisions and quotes with speakers. Hands off: Decision log.
  - Insight Analyst (skills: insight-generation + strategic-analysis): Finds the top insights, contradictions and weak signals; states confidence. Hands off: Insight table.
  - Briefer (skills: executive-briefing): Writes the two-page brief with owners and next actions. Hands off: Leadership brief.
4. STOP at Gate 1: After the Insight Analyst: the insights lead approves the top three insights, their confidence and the safety routing before the brief is written. Show me the evidence pack and wait.
5. Continue, then have an independent checker (deliverable-quality-review (reports to the human, not the lead)) review the draft. Fix what it finds.
6. STOP at Gate 2: After the Briefer: the head of medical signs off the brief, the owners and the actions before it is shared beyond the team. Wait for my decision.
7. Keep a timing note: minutes per step, and minutes of human review at each gate.

## House rules
- Never invent a citation, number, quote or meeting fact. Unsupported claims are marked and left for a human.
- Keep source IDs on every claim. Keep public evidence separate from fictional practice data.
- Do not send, post, publish or change any external system. Drafts only. Mark every output DRAFT and the data FICTIONAL.
```

All products, people and results in the practice data are fictional. Outputs are drafts that need qualified human review; this is not a clinical decision, GxP or pharmacovigilance system.

