



03

Congress Monitor

The congress just ended. What changed, and what do we do about it?

Hackathon vertical 03 of 04 · build afternoon agenda, swarm, gates, question bank and impact worksheet

01 Publications Brain

02 Insights Engine

03 Congress Monitor

04 Field Intelligence Engine

AI in Action for Medical Affairs

October 13–14, 2026 · Convene Philadelphia (2nd floor), Two Commerce Square, 2001 Market St

Host and keynote: Vivek Mukhatyar · aiinaction.up.railway.app/hackathon



Scan for the hackathon pack

The vertical on one page

THE CHALLENGE

Three days, sixty presentations, two competitor announcements, and leadership has four minutes. Today the readout lands a week later as a summary of what was presented, after people have already formed their impressions.

DAY 2 MORNING · WED OCT 14

3rd of 4 to present

01 02 **03** 04

About 20 min, including a few questions. Order 01 → 04, after the final build.

THE PROBLEM PROMPT

Build an AI workforce that, within hours of a congress closing, compares what was presented against what we expected and what our medical plan currently claims, applies the same scepticism to our data as to competitors', and tells leadership on one page what changed, why it matters, what to watch and what we should do, with every claim traced to an abstract or announcement. Humans approve the expectations and any recommended change.

Teams can keep this prompt or narrow it to one workflow they know. A good narrowed prompt names the role, the trigger, the deliverable and the decision it serves.

DONE LOOKS LIKE · DAY 1

The swarm ran end to end at least once on practice data; it stopped at Gate 1 and Gate 2 and the team made both decisions; screenshots and minute-level timings are in the capture log; the agent has drafted a slide outline.

DONE LOOKS LIKE · DAY 2

The deck is built from the final template: no [brackets] left, impact numbers are the team's own, practice data marked fictional, the LIVE DEMO slide left empty, backup screenshots in place. The team presents third of four, for about 20 minutes with a short live demo.

LEAD SKILL

congress-intelligence

SUPPORTING SKILLS

competitive-intelligence

evidence-appraisal

strategic-analysis

DATA PACK

Any TA. Each pack has congress-abstracts.md and competitor-announcements.md written for it.

congress-abstracts.md competitor-announcements.md

evidence-landscape.md medical-plan.md + practice CRM

Fictional practice data (Nordvant Biopharma), or bring your own non-confidential data.

WARM-UP MISSIONS

L2 congress

What changed after congress?

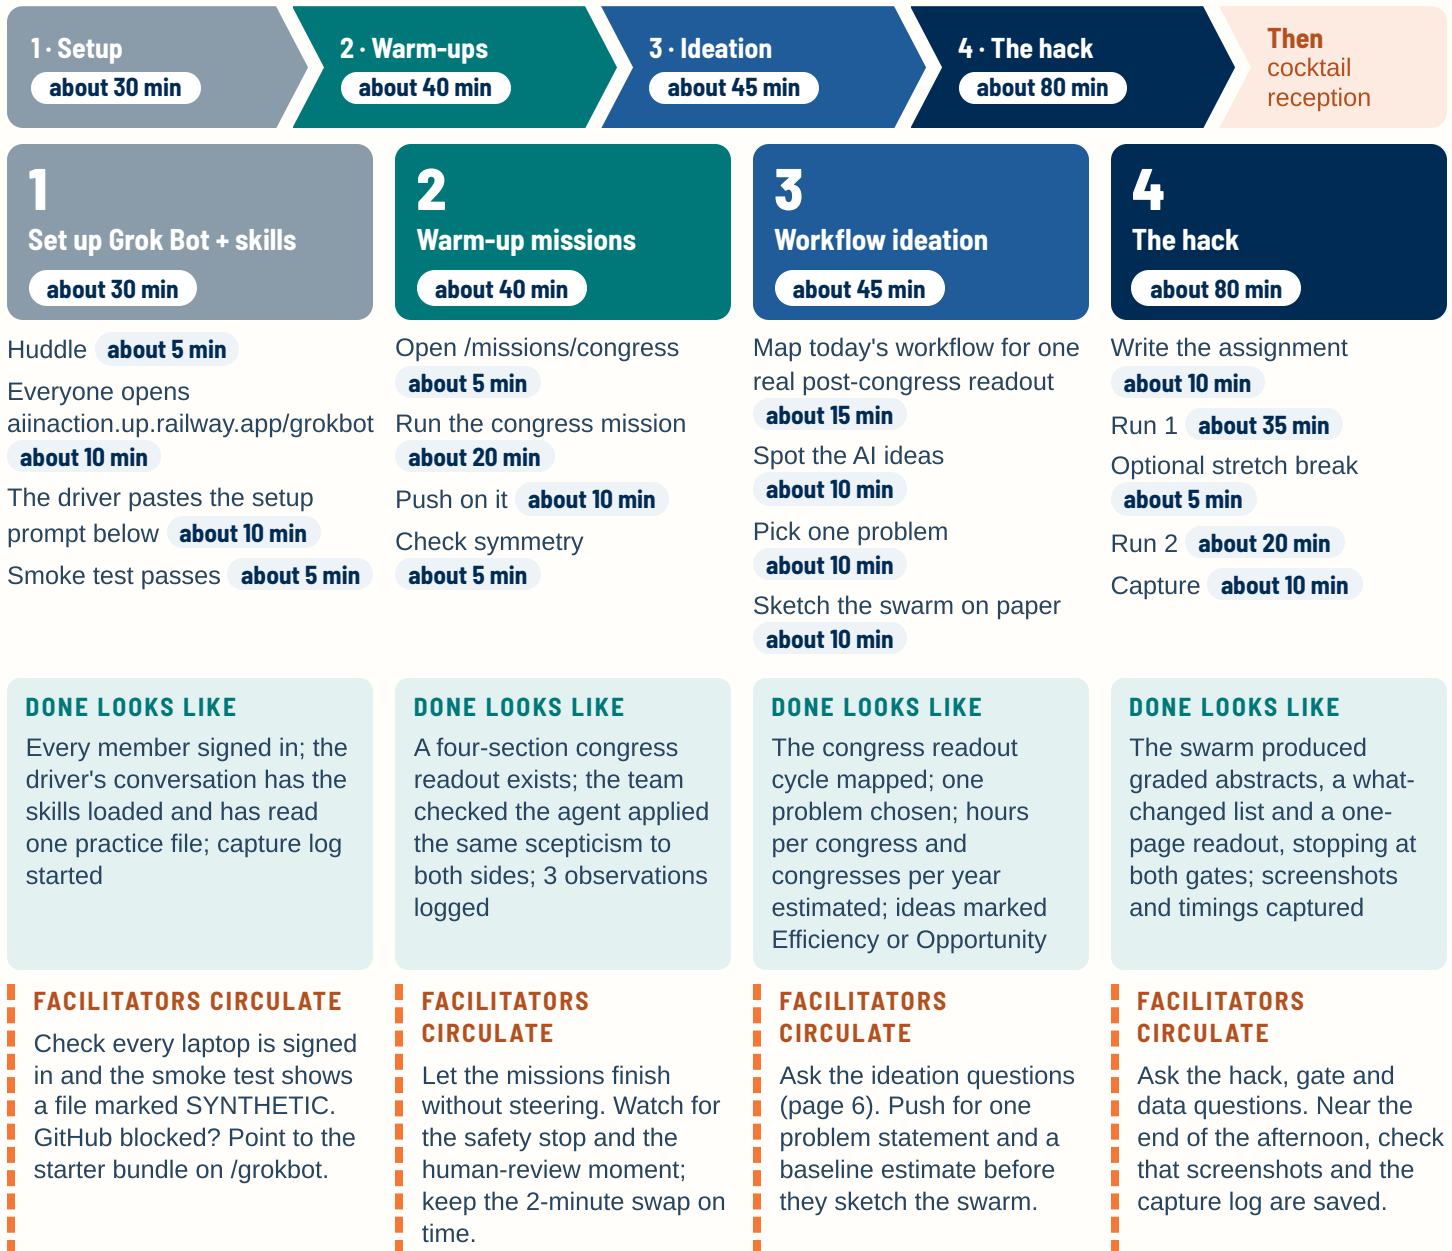
L1 kol-meeting

Prepare for the difficult meeting

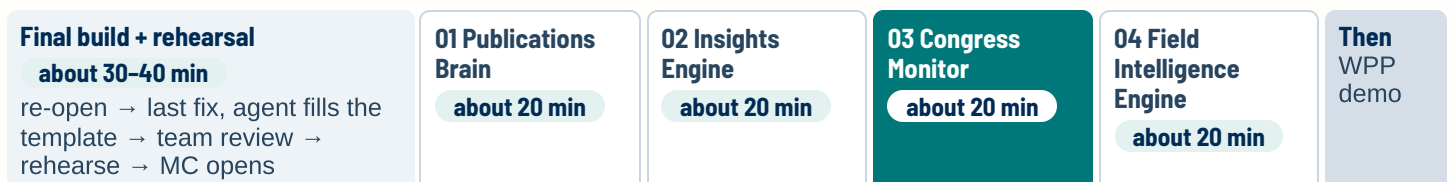


The build afternoon, block by block

A suggested, optional order of four blocks with approximate durations; teams may move faster or slower. **No formal share-out at the end of the afternoon:** facilitators circulate and ask questions (page 6).



DAY 2 · WED OCT 14 · FINAL BUILD → PRESENTATIONS IN ORDER 01 → 04



Congress Monitor presents third of four, about 20 minutes, give or take, including a short live demo and a few questions. Keep all decks on one laptop and treat the capture step near the end of Day 1 as essential.



Warm up, then pick your data

Warm-up missions Block 2 · about 40 min

Level 2 · Pair

Whole team (one driver)

"What changed after congress?"

The core task: congress findings compared with the existing medical plan; what changes, what does not, and why.

aiinaction.up.railway.app/missions/congress

Level 1 · Starter

Pair B (optional, in parallel)

"Prepare for the difficult meeting"

Prepare a post-congress conversation with a sceptical expert: what will she ask about the new data?

aiinaction.up.railway.app/missions/kol-meeting

STRETCH, ALTERNATIVES AND REFERENCES

L1 Starter

L2 Pair

L3 Team

L4 Swarm

L2 Alternative

What should leadership know?

</missions/field-insights>

L3 Hack reference

Work across a practice CRM and content library

</missions/connected-planning>

L3 Optional

Design an advisory board worth holding

</missions/advisory-board>

L4 Also useful

Build the whole medical launch plan

</missions/launch-plan-swarm>

BG Background

Team mission 3: the original team mission

[Medical-Affairs-Skills · workshop](#)

Data pack browse and copy links at aiinaction.up.railway.app/data

Fictional practice data, or bring your own non-confidential data.

Nordvant Biopharma and its products are fictional. Own data: no patient data, no confidential company data, no real HCP personal data.

PICK ONE TA

MM · NORVANTIB

AD · DERMALYX

Obesity · ADIPOSYN

Any TA. Each pack has congress-abstracts.md and competitor-announcements.md written for it.

TA PACK FILES synthetic/<TA>/

TA congress-abstracts.md

TA competitor-announcements.md

TA evidence-landscape.md

TA medical-plan.md

TA product-profile.md

TA guideline-landscape.md

TA kol-dossiers.md

TA field-observations.csv

PRACTICE CRM synthetic/connected/

CRM interactions.csv

CRM content_assets.csv

CRM hcps.csv

The practice CRM interactions table has a Congress channel: filter it to see what HCPs said at the congress. content_assets.csv shows which materials may now be out of date.



Efficiency vs Opportunity, and the digital team

Efficiency AI today's work, faster

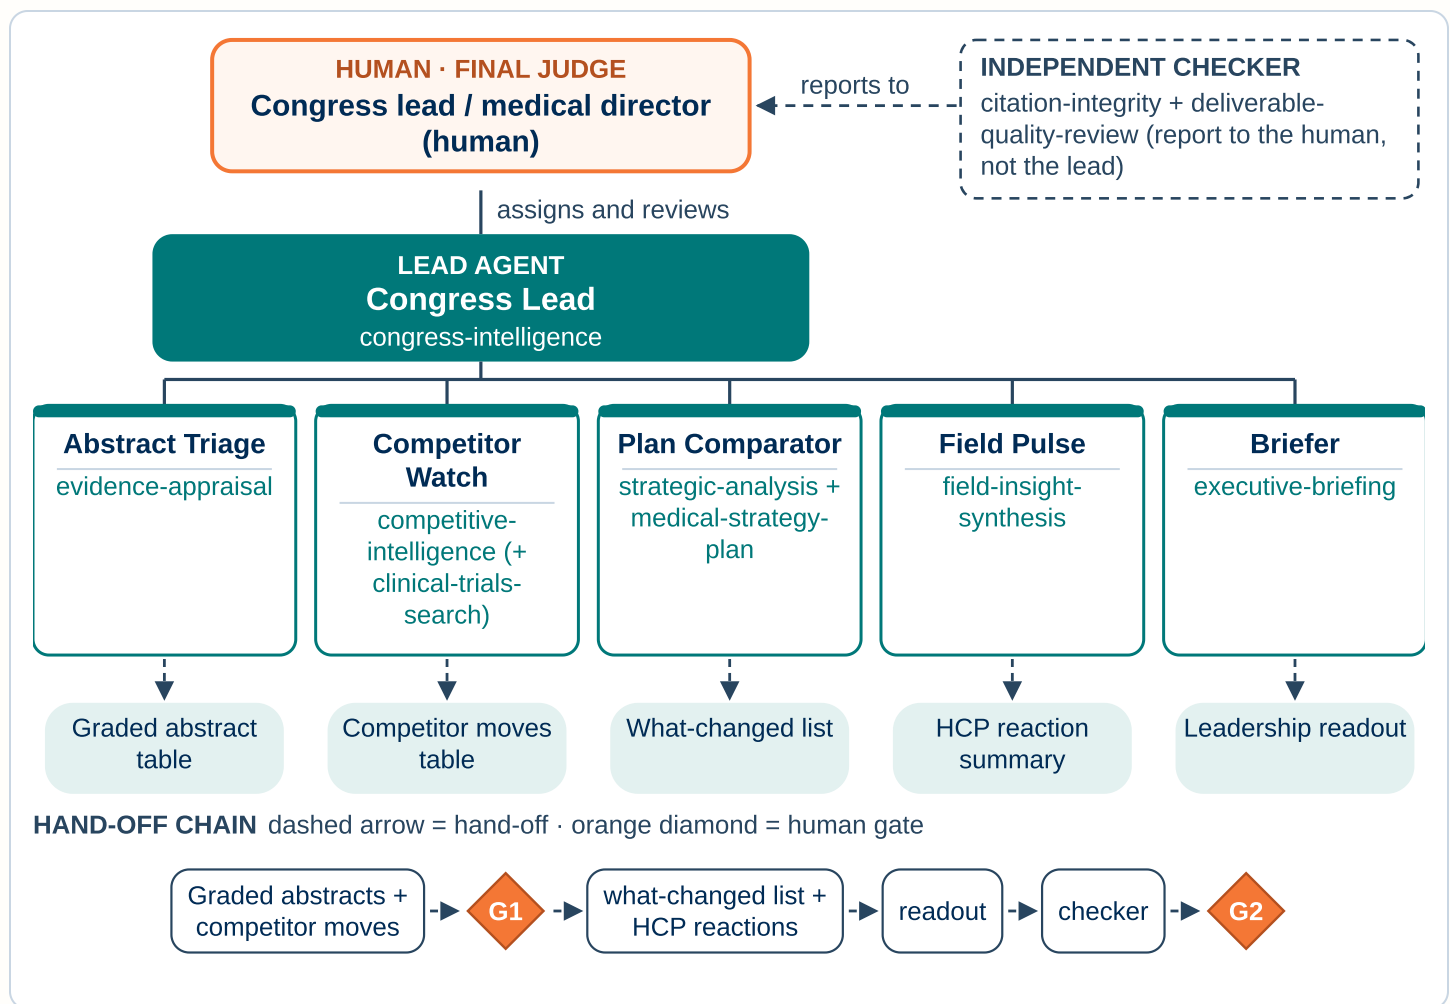
- E1** Abstract triage: every abstract tagged by relevance to our plan, with an evidence-quality grade, in the first hour after the congress.
- E2** Competitor announcement digest that separates what the data show from what the press release claims.
- E3** Same-day first draft of the four-section leadership readout (what changed, why it matters, what to watch, what to do).

vs

Opportunity AI something not possible today

- O1** Expectation-vs-reality diff: log our expectations before the congress, then get an automatic "where reality differed" analysis when it closes.
- O2** Cover every session and poster, not only the top ten, so weak signals and small studies are not missed.
- O3** Downstream impact in 24 hours: which of our claims, slides, FAQs and MSL materials are now outdated, plus draft field talking points for human approval.

Suggested swarm one lead agent, sub-agents with their skills, two human gates



Where humans decide, and what to ask



G1 Gate 1

WHEN	Before the Plan Comparator runs
WHO DECIDES	The congress lead
WHAT THEY CHECK	Approves the expectations register and the shortlist of findings that matter.
WHY A HUMAN	Deciding what "matters" for the strategy is judgment, and it stops the agent anchoring on the wrong comparison.

G2 Gate 2

WHEN	After the Briefer
WHO DECIDES	The medical director
WHAT THEY CHECK	Approves the recommendations before any plan, claim, material or field communication changes.
WHY A HUMAN	Changing scientific positions and field messaging needs accountable sign-off and compliance review.

Always-on stop (a rule, not a gate): any possible adverse event, product quality complaint or off-label signal in human-sourced records is listed verbatim and routed by a human before analysis continues.

Facilitator question bank the Day 1 feedback loop: ask one or two per visit, then move on

Ideation block 3

- 1 How long after the congress does leadership get the readout today? What decisions are made before it arrives?
- 2 What did your team expect before this congress? Where is that written down today?
- 3 Who decides which sessions to cover today, and what gets missed?
- 4 Which downstream materials (slides, FAQs, MSL decks) usually go out of date after a congress?

The hack block 4

- 5 Is the swarm as sceptical about our data as about the competitor's? Show me one example of each.
- 6 What does the agent do with a single-arm study or a press-release claim with no data?
- 7 Which recommendation would change something external, and who signs that off?

Gates during the hack

- 8 Where is Gate 1, and what does the congress lead approve there?

Data during the hack

- 9 Show me the "expected vs what happened" comparison. Where did reality differ?

Impact during the hack

- 10 How many people-hours does a congress readout take today, from session coverage to leadership deck? Per congress, how many congresses a year?



When a team gets stuck

The ten snags facilitators see most often in this vertical, each with the unblocker to offer. Nudge, then step back: the team does the building.

Pitfall	How to unblock
01 A summary of what was presented, not analysis.	Ask for the comparison against prior expectations and the current plan; load strategic-analysis.
02 Asymmetric scepticism (tough on them, generous on us).	Give the same evidence-appraisal checklist to both; the checker compares the two.
03 Abstracts cited as if they were published evidence.	Mark abstract-level data as preliminary; citation-integrity flags every overreach.
04 Press-release language copied into the readout.	Competitor Watch separates "what the data show" from "what they claim".
05 Sixty abstracts is a lot for 80 minutes.	Start with the ten most relevant (the triage worker picks, a human confirms), then widen.
06 Talking points drift into promotion.	Field talking points are drafts for medical and compliance review only (Gate 2).
07 Grok Bot cannot open GitHub.	Upload the starter bundle for your TA from https://aiinaction.up.railway.app/grokbob, or the repository ZIP.
08 The agent asks to connect company systems.	Say no: the workshop uses practice data. It is in every setup prompt.
09 A run is slow or wanders.	Ask for the plan first, then run one sub-agent at a time; cut the scope to one TA, one HCP, one output.
10 The output is generic.	Teach the agent: add 3–5 house rules (what good looks like, what must never appear) and run again.



Unblock, then step back. Teams learn most when they make the call at each gate themselves.



Impact worksheet and setup prompt

Impact worksheet the team's own estimates

One cycle = one major congress readout (coverage, analysis, leadership deck, field update)

Today baseline	h	×		×	\$ /h	=	\$
	hours per cycle (all people)				loaded rate		cost per year
vs							
AI-assisted with the swarm	h	×		×	\$ /h	=	\$
	hours per cycle, incl. human review				loaded rate		cost per year

HOURS SAVED PER YEAR (today h – AI h) × cycles =	VALUE PER YEAR (ESTIMATE) hours saved × rate =	OTHER BENEFITS
h	\$	days % # speed to decision coverage errors caught

ILLUSTRATIVE

Not data, only an example: 60 h → 15 h per cycle × 4 cycles/year = 180 h saved; × \$150/h = \$27,000 per year. Replace every number with the team's own estimate.

Count every person involved and include human review time in the AI-assisted row. These numbers go on the Impact slide of the final template (slide 9: formula tiles and a bar chart built from the team's own estimates).

Setup prompt Block 1, after sign-in · the driver pastes it into Grok Bot (replace the TA)

1 · Setup prompt

copy all

Use <https://github.com/Open-Medical-Affairs/Medical-Affairs-Skills>
 Read AGENTS.md. We are the Congress Monitor team at AI in Action for Medical Affairs.
 Load medical-affairs-foundations and these skills: congress-intelligence, competitive-intelligence, evidence-appraisal, strategic-analysis, executive-briefing, citation-integrity, deliverable-quality-review.
 Use the practice data in <https://github.com/Open-Medical-Affairs/Data-Sources>: synthetic/[oncology-mm | immunology-ad | cardiometabolic-obesity]/ and synthetic/connected/ (the practice CRM). It is fictional (Nordvant Biopharma).
 First, only: list the skills you loaded, then show the first 5 rows of interactions.csv and confirm the file is labelled SYNTHETIC. Do not start a mission yet.
 Do not ask me to connect company systems. Do not send, post or change anything outside this conversation.

SUGGESTED ROLES (ROTATE IF YOU LIKE)

Driver

types into Grok Bot

Navigator

reads the mission and pushes back

Scribe

owns the capture log

Timekeeper + presenter

keeps the clock, tells the story



Assignment starter for the hack

Paste into Grok Bot at the start of Block 4 (the hack), then edit the [brackets]. The setup prompt for Block 1 is on page 8. Ready prompts for every mission: aiinaction.up.railway.app/prompts · swarm mode and the Skill creator: [/optimizer](#).

2 · Assignment starter: Congress Monitor swarm

Block 4 · start of the hack

```
# Assignment: Congress Monitor swarm
You are the Congress Lead, a lead agent coordinating a small team of digital workers for a Medical Affairs team. This is an assignment, not a question: plan, delegate, check and hand back finished drafts.

## Goal: what done looks like
[One sentence: the problem we picked in ideation, and the deliverable that solves it.]

## Context to use
- https://github.com/Open-Medical-Affairs/Medical-Affairs-Skills : read AGENTS.md, load medical-affairs-foundations, your lead skill congress-intelligence, and the skills named below.
- Practice data (fictional): https://github.com/Open-Medical-Affairs/Data-Sources synthetic/[TA]/ and synthetic/connected/. Or our own non-confidential data: [describe].

## How to work
1. Show me a short plan and the org chart of workers first. Wait for my OK.
2. Scan human-sourced records for possible safety findings (adverse events, product complaints, off-label use) before anything else. List them verbatim and stop until I acknowledge.
3. Run these sub-agents, each with only the context it needs:
  - Abstract Triage (skills: evidence-appraisal): Reads and grades every abstract; tags relevance to our plan. Hands off: Graded abstract table.
  - Competitor Watch (skills: competitive-intelligence (+ clinical-trials-search)): Separates data from spin in competitor announcements. Hands off: Competitor moves table.
  - Plan Comparator (skills: strategic-analysis + medical-strategy-plan): Compares findings with our expectations and current claims. Hands off: What-changed list.
  - Field Pulse (skills: field-insight-synthesis): Reads Congress-channel CRM interactions for HCP reactions. Hands off: HCP reaction summary.
  - Briefer (skills: executive-briefing): Writes the one-page, four-section readout. Hands off: Leadership readout.
4. STOP at Gate 1: Before the Plan Comparator runs: the congress lead approves the expectations register and the shortlist of findings that matter. Show me the evidence pack and wait.
5. Continue, then have an independent checker (citation-integrity + deliverable-quality-review (report to the human, not the lead)) review the draft. Fix what it finds.
6. STOP at Gate 2: After the Briefer: the medical director approves the recommendations before any plan, claim, material or field communication changes. Wait for my decision.
7. Keep a timing note: minutes per step, and minutes of human review at each gate.

## House rules
- Never invent a citation, number, quote or meeting fact. Unsupported claims are marked and left for a human.
- Keep source IDs on every claim. Keep public evidence separate from fictional practice data.
- Do not send, post, publish or change any external system. Drafts only. Mark every output DRAFT and the data FICTIONAL.
```

All products, people and results in the practice data are fictional. Outputs are drafts that need qualified human review; this is not a clinical decision, GxP or pharmacovigilance system.

