



04

Field Intelligence Engine

Walk into every HCP conversation prepared, and close the loop after.

Hackathon vertical 04 of 04 · build afternoon agenda, swarm, gates, question bank and impact worksheet

01 Publications Brain

02 Insights Engine

03 Congress Monitor

04 Field Intelligence Engine

AI in Action for Medical Affairs

October 13–14, 2026 · Convene Philadelphia (2nd floor), Two Commerce Square, 2001 Market St

Host and keynote: Vivek Mukhatyar · aiinaction.up.railway.app/hackathon



Scan for the hackathon pack

The vertical on one page

THE CHALLENGE

MSLs spend hours before each scientific exchange stitching together CRM history, open enquiries and evidence, and hours afterwards writing CRM notes and follow-ups. Preparation quality varies, and commitments made in the meeting slip.

DAY 2 SLOT · WED OCT 14

10:40–11:00 AM

About 20 minutes, give or take, including a few questions. Teams present 01 → 04 from 9:40.

THE PROBLEM PROMPT

Build an AI workforce that, for any HCP on an MSL's list, prepares a one-page pre-call brief from verified professional context, prior interactions, open questions and permitted materials; and after the call turns the notes into a factual CRM draft, a commitment register and a follow-up draft, routing any safety finding or unsolicited off-label request to the right human. The MSL approves everything; nothing is sent or logged by the agent.

Teams can keep this prompt or narrow it to one workflow they know. A good narrowed prompt names the role, the trigger, the deliverable and the decision it serves.

DONE LOOKS LIKE · DAY 1

The swarm ran end to end at least once on practice data; it stopped at Gate 1 and Gate 2 and the team made both decisions; screenshots and minute-level timings are in the capture log; the agent has drafted a slide outline.

DONE LOOKS LIKE · DAY 2

The deck is built from the final template: no [brackets] left, impact numbers are the team's own, practice data marked fictional, the LIVE DEMO slide left empty, backup screenshots in place. The team presents at 10:40 for about 20 minutes with a short live demo.

LEAD SKILL

field-medical-planning

SUPPORTING SKILLS

msl-pre-call-planning

kol-engagement-brief

hcp-discovery-and-access

DATA PACK

Any TA. The practice CRM (connected) spans all three TAs; pick HCPs from one TA to keep it focused.

kol-dossiers.md interaction-notes.md

field-account-plan.csv product-profile.md

evidence-landscape.md + practice CRM

Fictional practice data (Nordvant Biopharma), or bring your own non-confidential data.

WARM-UP MISSIONS

L1 msl-pre-call

Prepare the MSL pre-call brief

L1 msl-post-call

Turn the call notes into useful follow-up



The build afternoon, block by block

A suggested, optional flow: **Grok Bot setup (30 min)** → **warm-up missions** → **workflow ideation** → **the hack**. Teams may move faster or slower. **5:00 PM: no formal share-out**; facilitators circulate and ask questions (page 6), then the cocktail reception.

1 · Setup 30 min	2 · Warm-ups 40 min	3 · Ideation 45 min	4 · The hack 80 min
1:45	2:15	2:55	3:40
5:00			
1 Set up Grok Bot + skills 1:45–2:15 · 30 min	2 Warm-up missions 2:15–2:55 · 40 min	3 Workflow ideation 2:55–3:40 · 45 min	4 The hack 3:40–5:00 · 80 min
1:45 Huddle 1:50 Everyone opens <code>aiinaction.up.railway.app/grokbob</code> 2:02 The driver pastes the setup prompt below 2:10 Smoke test passes	2:15 Split into two pairs 2:18 Run both Level 1 missions in parallel 2:36 Push on it 2:48 Connect the two	2:55 Map today's workflow for one real HCP engagement 3:10 Spot the AI ideas 3:22 Pick one problem 3:32 Sketch the swarm on paper	3:40 Write the assignment 3:50 Run 1 4:25 Optional stretch break 4:30 Run 2 4:50 Capture
DONE LOOKS LIKE Every member signed in; the driver's conversation has the skills loaded and has read one practice file; capture log started	DONE LOOKS LIKE Pair A has a pre-call brief; Pair B has a CRM draft and commitment register; the team saw how one could feed the other; 3 observations logged	DONE LOOKS LIKE One HCP engagement mapped end to end; one problem chosen; prep and write-up hours per call and calls per week estimated; ideas marked Efficiency or Opportunity	DONE LOOKS LIKE The swarm ran one HCP end to end (brief, notes, CRM draft, follow-up), stopping at both gates; screenshots and timings captured
FACILITATORS CIRCULATE Check every laptop is signed in and the smoke test shows a file marked SYNTHETIC. GitHub blocked? Point to the starter bundle on /grokbob.	FACILITATORS CIRCULATE Let the missions finish without steering. Watch for the safety stop and the human-review moment; keep the 2-minute swap on time.	FACILITATORS CIRCULATE Ask the ideation questions (page 6). Push for one problem statement and a baseline estimate before they sketch the swarm.	FACILITATORS CIRCULATE Ask the hack, gate and data questions. From 4:45, check that screenshots and the capture log are saved.

DAY 2 · WED OCT 14 · FINAL BUILD → THIS TEAM'S SLOT

9:00–9:38 · Final build + rehearsal re-open → last fix, agent fills the template → team review → rehearse · MC opens 9:38	9:40 01 Publications Brain	10:00 02 Insights Engine	10:20 03 Congress Monitor	10:40 04 Field Intelligence Engine	11:00 WPP demo
---	--------------------------------------	------------------------------------	-------------------------------------	--	--------------------------

Field Intelligence Engine: 10:40–11:00 AM, about 20 minutes, give or take, including a short live demo and a few questions. Teams present 01 → 04 from 9:40. Keep all decks on one laptop and treat the Day 1 4:50–5:00 capture as essential.



Warm up, then pick your data

Warm-up missions Block 2 · 2:15–2:55

Level 1 · Starter

Pair A

“Prepare the MSL pre-call brief”

Before the call: a one-page brief from CRM history, open tasks and access context, with no invented meeting facts.

aiinaction.up.railway.app/missions/msl-pre-call

Level 1 · Starter

Pair B

“Turn the call notes into useful follow-up”

After the call: notes to a CRM draft, a commitment and gap register, and a follow-up draft.

aiinaction.up.railway.app/missions/msl-post-call

STRETCH, ALTERNATIVES AND REFERENCES

L1 Starter

L2 Pair

L3 Team

L4 Swarm

L1 Alternative

Clear the administrative queue
</missions/msl-admin>

L1 Alternative

Prepare for the difficult meeting
</missions/kol-meeting>

L2 Alternative

Identify scientific coverage and access gaps
</missions/hcp-access>

L3 Hack reference

Work across a practice CRM and content library
</missions/connected-planning>

L4 Also useful

Build the whole medical launch plan
</missions/launch-plan-swarm>

BG Background

Team mission 1: the original team mission
[Medical-Affairs-Skills · workshop](#)

Data pack browse and copy links at aiinaction.up.railway.app/data

Fictional practice data, or bring your own non-confidential data.

Nordvant Biopharma and its products are fictional. Own data: no patient data, no confidential company data, no real HCP personal data.

PICK ONE TA

MM · NORVANTIB

AD · DERMALYX

Obesity · ADIPOSYN

Any TA. The practice CRM (connected) spans all three TAs; pick HCPs from one TA to keep it focused.

TA PACK FILES synthetic/<TA>/

TA kol-dossiers.md

TA interaction-notes.md

TA field-account-plan.csv

TA product-profile.md

TA evidence-landscape.md

TA medical-plan.md

PRACTICE CRM synthetic/connected/

CRM hcps.csv

CRM interactions.csv

CRM accounts.csv

CRM access_log.csv

CRM msl_tasks.csv

CRM enquiries.csv

CRM content_assets.csv

CRM data-dictionary.json

The practice CRM also ships as one SQLite file ([medical-affairs.sqlite](#)) if the agent prefers to query it. All HCPs, accounts and MSLs are fictional; never paste real HCP names.



Efficiency vs Opportunity, and the digital team

Efficiency AI today's work, faster


- E1** Pre-call brief in five minutes: prior interactions, open enquiries, open tasks, permitted materials and suggested questions on one page.
- E2** Post-call CRM note, commitment register and follow-up draft from the MSL's notes, ready for the MSL to verify.
- E3** Weekly admin queue cleared: overdue tasks, missing records and the status brief assembled from the CRM.

vs

Opportunity AI something not possible today

- O1** A brief before every HCP interaction, not only the top KOLs: preparation quality becomes the same across the whole team.
- O2** Coverage-gap radar: HCPs with unanswered scientific questions, or no scientific contact this quarter, surfaced weekly with an ethical access route.
- O3** Closed loop: every call automatically feeds insights to the Insights team and unsolicited questions to Medical Information, so nothing said in the field is lost.

Suggested swarm one lead agent, sub-agents with their skills, two human gates



Where humans decide, and what to ask



G1

Gate 1

WHEN	Before the call
WHO DECIDES	The MSL
WHAT THEY CHECK	Approves the brief and the planned questions (non-promotional, scientifically balanced, within permitted materials).
WHY A HUMAN	The MSL is accountable for the scientific exchange and for compliance.

G2

Gate 2

WHEN	After the call
WHO DECIDES	The MSL
WHAT THEY CHECK	Verifies the CRM note, the commitments and the follow-up drafts before anything is logged or sent; safety findings and off-label requests go to the right function.
WHY A HUMAN	CRM records are official records and correspondence leaves the company.

Always-on stop (a rule, not a gate): any possible adverse event, product quality complaint or off-label signal in human-sourced records is listed verbatim and routed by a human before analysis continues.

Facilitator question bank the Day 1 feedback loop: ask one or two per visit, then move on

Ideation 2:55–3:40

- 1 How long does an MSL spend preparing for one scientific exchange today? For which HCPs do they skip it?
- 2 Where does the information for a brief live today, and which source is least trusted?
- 3 What happens today to a commitment an MSL makes in a meeting? How often does it slip?
- 4 Which part of the call itself must always stay human?

The hack 3:40–5:00

- 5 Is anything in the brief promotional or off-label? How would the swarm catch it?

Gates during the hack

- 6 What happens when the call notes contain a possible adverse event? Who gets it, and how fast?
- 7 Where are Gate 1 and Gate 2, and what does the MSL see at each?

Data during the hack

- 8 Show me a fact in the brief and the CRM row it came from. What did the agent refuse to guess?
- 9 What personal information about the HCP is the agent allowed to use? Professional context only?

Impact during the hack

- 10 How many calls per MSL per week? How long do prep and write-up take today versus with the swarm, including the MSL's review?



When a team gets stuck

The ten snags facilitators see most often in this vertical, each with the unblocker to offer. Nudge, then step back: the team does the building.

Pitfall	How to unblock
01 The agent invents meeting facts or HCP details.	● Rule: every fact in the brief cites a CRM row; unknowns go to "open questions".
02 Promotional tone in briefs or follow-ups.	● Load kol-engagement-brief and the checker; Gate 1 includes a non-promotional check.
03 Sending or logging without review.	● The agent drafts only. "Do not send, post or change any system" goes in the assignment.
04 Real HCP names typed in from memory.	● Practice data only, or your own non-confidential, de-identified data.
05 Trying to cover the whole territory at once.	● Start with one HCP end to end (brief, call notes, follow-up), then scale to five.
06 Access data used unethically.	● hcp-discovery-and-access proposes institution-aware routes; humans decide on any outreach.
07 Grok Bot cannot open GitHub.	● Upload the starter bundle for your TA from https://aiinaction.up.railway.app/grokbob, or the repository ZIP.
08 The agent asks to connect company systems.	● Say no: the workshop uses practice data. It is in every setup prompt.
09 A run is slow or wanders.	● Ask for the plan first, then run one sub-agent at a time; cut the scope to one TA, one HCP, one output.
10 The output is generic.	● Teach the agent: add 3–5 house rules (what good looks like, what must never appear) and run again.



Unblock, then step back. Teams learn most when they make the call at each gate themselves.



Impact worksheet and setup prompt

Impact worksheet the team's own estimates

One cycle = one HCP call (prep + CRM note + follow-up), per MSL

Today baseline	h hours per cycle (all people)	×	 cycles per year	×	\$ /h loaded rate	=	\$ cost per year
vs							
AI-assisted with the swarm	h hours per cycle, incl. human review	×	 cycles per year	×	\$ /h loaded rate	=	\$ cost per year

HOURS SAVED PER YEAR (today h – AI h) × cycles = h	VALUE PER YEAR (ESTIMATE) hours saved × rate = \$	OTHER BENEFITS
		days speed to decision % coverage # errors caught

ILLUSTRATIVE

Not data, only an example: 2.5 h → 0.75 h per cycle × 200 cycles/year = 350 h saved; × \$120/h = \$42,000 per year. Replace every number with the team's own estimate.

Count every person involved and include human review time in the AI-assisted row. These numbers go on the Impact slide of the final template (slide 9: formula tiles and a bar chart built from the team's own estimates).

Setup prompt Block 1 · 2:02 · the driver pastes it into Grok Bot (replace the TA)

1 · Setup prompt

[copy all](#)

Use <https://github.com/Open-Medical-Affairs/Medical-Affairs-Skills>
 Read AGENTS.md. We are the Field Intelligence Engine team at AI in Action for Medical Affairs.
 Load medical-affairs-foundations and these skills: field-medical-planning, msl-pre-call-planning, kol-engagement-brief, msl-post-call-follow-up, hcp-discovery-and-access, msl-administrative-operations, deliverable-quality-review.
 Use the practice data in <https://github.com/Open-Medical-Affairs/Data-Sources>: synthetic/[oncology-mm | immunology-ad | cardiometabolic-obesity]/ and synthetic/connected/ (the practice CRM). It is fictional (Nordvant Biopharma).
 First, only: list the skills you loaded, then show the first 5 rows of hcps.csv and confirm the file is labelled SYNTHETIC. Do not start a mission yet.
 Do not ask me to connect company systems. Do not send, post or change anything outside this conversation.

SUGGESTED ROLES (ROTATE IF YOU LIKE)

Driver

types into Grok Bot

Navigator

reads the mission and pushes back

Scribe

owns the capture log

Timekeeper + presenter

keeps the clock, tells the story



Assignment starter for the hack

Paste into Grok Bot at the start of Block 4 (3:40), then edit the [brackets]. The setup prompt for Block 1 is on page 8. Ready prompts for every mission: aiinaction.up.railway.app/prompts · swarm mode and the Skill creator: [/optimizer](#).

2 · Assignment starter: Field Intelligence Engine swarm

Block 4 · 3:40

```
# Assignment: Field Intelligence Engine swarm
You are the Field Copilot, a lead agent coordinating a small team of digital workers for a Medical Affairs team. This is an assignment, not a question: plan, delegate, check and hand back finished drafts.

## Goal: what done looks like
[One sentence: the problem we picked in ideation, and the deliverable that solves it.]

## Context to use
- https://github.com/Open-Medical-Affairs/Medical-Affairs-Skills : read AGENTS.md, load medical-affairs-foundations, your lead skill field-medical-planning, and the skills named below.
- Practice data (fictional): https://github.com/Open-Medical-Affairs/Data-Sources synthetic/[TA]/ and synthetic/connected/. Or our own non-confidential data: [describe].

## How to work
1. Show me a short plan and the org chart of workers first. Wait for my OK.
2. Scan human-sourced records for possible safety findings (adverse events, product complaints, off-label use) before anything else. List them verbatim and stop until I acknowledge.
3. Run these sub-agents, each with only the context it needs:
  - Pre-call Briefer (skills: msl-pre-call-planning + kol-engagement-brief): Builds the one-page brief from CRM history and permitted materials. Hands off: Pre-call brief.
  - Coverage Mapper (skills: hcp-discovery-and-access): Finds coverage gaps and ethical access routes. Hands off: Coverage register.
  - Post-call Scribe (skills: msl-post-call-follow-up): Turns call notes into a CRM draft and a commitment register. Hands off: CRM draft + commitments.
  - Follow-up Drafter (skills: medical-information-response + medical-correspondence): Drafts answers to unsolicited questions and the follow-up letter. Hands off: Follow-up drafts.
  - Admin Clerk (skills: msl-administrative-operations): Updates the task queue and the weekly status brief. Hands off: Task queue.
4. STOP at Gate 1: Before the call: the MSL approves the brief and the planned questions (non-promotional, scientifically balanced, within permitted materials). Show me the evidence pack and wait.
5. Continue, then have an independent checker (deliverable-quality-review (reports to the MSL, not the lead)) review the draft. Fix what it finds.
6. STOP at Gate 2: After the call: the MSL verifies the CRM note, the commitments and the follow-up drafts before anything is logged or sent; safety findings and off-label requests go to the right function. Wait for my decision.
7. Keep a timing note: minutes per step, and minutes of human review at each gate.

## House rules
- Never invent a citation, number, quote or meeting fact. Unsupported claims are marked and left for a human.
- Keep source IDs on every claim. Keep public evidence separate from fictional practice data.
- Do not send, post, publish or change any external system. Drafts only. Mark every output DRAFT and the data FICTIONAL.
```

All products, people and results in the practice data are fictional. Outputs are drafts that need qualified human review; this is not a clinical decision, GxP or pharmacovigilance system.

